

Phase 2 GMM Final

Jessala Grijalva

2024-03-20

Introduction

Setup and Preliminaries

Establishes the environment, loads necessary libraries, and sets the seed for reproducibility. Loads and prepares data for analysis, which is a crucial first step in any data analysis process.

Libraries and Seed

```
# Load necessary libraries
library(pacman)
p_load(mclust, cluster, dplyr, tidyverse, factoextra, flexclust, caret,
       dunn.test, MASS, Rtsne, NbClust, boot, fpc, ggplot2, gridExtra, grid,
       kableExtra, rstatix, PMCMRplus, FSA, DescTools, mclust, stats, purrr,
       fmsb, GGally, tinytex)

# Set seed for reproducibility
set.seed(2500)
```

Data Loading and Preparation

```
# Set the working directory to the folder where the RDA file is located
load("~/Desktop/Dissertation Analysis/Dissertation Analysis/Data/Processed/lns_clean.rda")

# Check for missing data in the dataset
missing_data_summary <- sapply(lns_subset, function(x) sum(is.na(x)))

# Print the summary of missing data
print(missing_data_summary)
```

## CULTURAL_IDENTITY	AMERICAN	LEARNENG	DISTINCT
## 0	0	0	0
## KEEPSPLAN	BLEND	VOTE_PREF	IDEOLOGY
## 0	0	0	0
## FEELPART	PARTYID	SAYSO	GOVTRUST
## 0	0	0	0
## INCSUPP	HEALTH	DREAMACT	IMMVIEW
## 0	0	0	0
## IMMPOLICY	EDUCATION	NATIONALITY	BIRTH_ORIGIN
## 0	0	0	0
## GEN_STATUS	AGE	SEX	RACE1
## 0	0	0	0

```

##          HHINC        RELIGION        WEIGHT
##          0             0             0

```

Clustering Preparation

Defines variables for clustering and standardizes them, which is necessary to ensure that all variables contribute equally to the analysis.

Define Clustering Variables

```

# Select variables for K Means clustering
selected_vars <- c("AMERICAN", "CULTURAL_IDENTITY", "KEEPSPAN", "DISTINCT",
                  "LEARNENG")

```

Standardization

```

# Standardize the selected variables. Directly converting the result into a matrix format suitable for clustering
clustering_data_matrix <- scale(lns_subset[selected_vars])

# Check if the standardization resulted in a matrix and print the dimensions to confirm
print(dim(clustering_data_matrix))

## [1] 4785      5

```

Clustering Evaluation Functions

Provides functions to calculate key clustering metrics. This is important to assess the quality of the clustering solution.

```

# Define the range of cluster numbers to evaluate
k_values <- c(3, 4, 5)

# Initialize a list to store GMM results and evaluation metrics for each k
gmm_results <- list()

for (k in k_values) {
  # Perform GMM clustering
  gmm_model <- Mclust(clustering_data_matrix, G = k)

  # Store the model
  gmm_results[[paste("GMM_k", k)]] <- gmm_model

  # Print basic model information including BIC
  cat("GMM Model for k =", k, "\n")
  print(summary(gmm_model, parameters = FALSE))
}

## GMM Model for k = 3
## -----
## Gaussian finite mixture model fitted by EM algorithm
## -----
## 
## Mclust EEV (ellipsoidal, equal volume and shape) model with 3 components:
## 
##   log-likelihood    n df      BIC      ICL

```

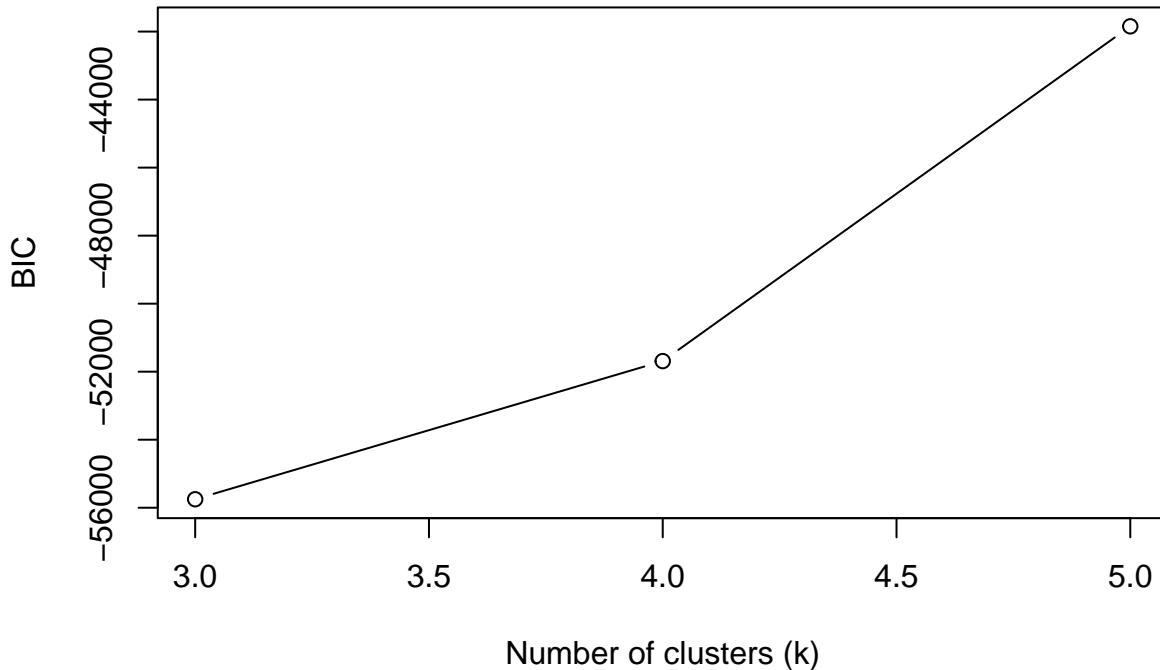
```

##      -27654.76 4785 52 -55750.13 -57658.22
##
## Clustering table:
##   1   2   3
## 724 2600 1461
## GMM Model for k = 4
## -----
## Gaussian finite mixture model fitted by EM algorithm
## -----
## 
## Mclust EEV (ellipsoidal, equal volume and shape) model with 4 components:
## 
## log-likelihood    n df      BIC      ICL
##          -25556.28 4785 68 -51688.75 -52293.62
##
## Clustering table:
##   1   2   3   4
## 433 705 378 3269
## GMM Model for k = 5
## -----
## Gaussian finite mixture model fitted by EM algorithm
## -----
## 
## Mclust EEV (ellipsoidal, equal volume and shape) model with 5 components:
## 
## log-likelihood    n df      BIC      ICL
##          -20567.52 4785 84 -41846.8 -42774.64
##
## Clustering table:
##   1   2   3   4   5
## 275 106 3135 604 665

# Extract and plot BIC scores from the GMM results
bic_values <- sapply(gmm_results, function(x) x$bic)
plot(k_values, bic_values, type = "b", xlab = "Number of clusters (k)", ylab = "BIC", main = "BIC Scores")

```

BIC Scores for Different k Values



GMM Means Summary

```
# Convert the original selected variables in the dataset to a matrix format for compatibility
original_data_matrix <- as.matrix(lns_subset[selected_vars])

# Initialize a list to store the mean values for each k
cluster_means_list <- list()

for (k in k_values) {
  # Retrieve the model
  gmm_model <- gmm_results[[paste("GMM_k", k)]]

  # Retrieve the cluster assignment from the model
  cluster_assignment <- gmm_model$classification

  # Calculate means for each cluster
  cluster_means <- aggregate(original_data_matrix, by = list(cluster = cluster_assignment), FUN = mean)

  # Store the calculated means
  cluster_means_list[[paste("Means_k", k)]] <- cluster_means

  # Print the means for each cluster for the current k
  cat("Cluster Means for k =", k, ":\n")
  print(cluster_means)
}

## Cluster Means for k = 3 :
##   cluster AMERICAN CULTURAL_IDENTITY KEEPSpan DISTINCT LEARNENG
## 1      1 3.473757      3.183011 3.261050 2.406077 3.243094
```

```

## 2      2 3.737692      3.428846 3.981154 2.701154 3.997308
## 3      3 2.657084      3.591718 3.617385 2.981520 3.995209
## Cluster Means for k = 4 :
##   cluster AMERICAN CULTURAL_IDENTITY KEEPSpan DISTINCT LEARNENG
## 1      1 3.233256      3.411085 3.635104 2.711316 2.766744
## 2      2 3.880851      3.052482 3.327660 1.797163 3.988652
## 3      3 3.523810      3.158730 2.677249 2.997354 3.947090
## 4      4 3.256959      3.561946 3.996635 2.920465 4.000000
## Cluster Means for k = 5 :
##   cluster AMERICAN CULTURAL_IDENTITY KEEPSpan DISTINCT LEARNENG
## 1      1 3.149091      3.643636 4.000000 2.767273 2.752727
## 2      2 3.481132      2.858491 3.207547 2.339623 2.764151
## 3      3 3.301435      3.575279 3.955343 2.999681 3.971930
## 4      4 3.344371      3.217715 3.973510 1.847682 4.000000
## 5      5 3.774436      3.022556 2.642105 2.393985 4.000000
# Convert the original selected variables in the dataset to a matrix format for compatibility
original_data_matrix <- as.matrix(lns_subset[selected_vars])

# Initialize a list to store the mean values and sizes for each k
cluster_means_and_sizes_list <- list()

for (k in k_values) {
  # Retrieve the model
  gmm_model <- gmm_results[[paste("GMM_k", k)]]

  # Retrieve the cluster assignment from the model
  cluster_assignment <- gmm_model$classification

  # Calculate means for each cluster
  cluster_means <- aggregate(original_data_matrix, by = list(cluster = cluster_assignment), FUN = mean)

  # Calculate cluster sizes
  cluster_sizes <- table(cluster_assignment)

  # Convert cluster_sizes to a dataframe for easy merging
  cluster_sizes_df <- as.data.frame(cluster_sizes)
  names(cluster_sizes_df) <- c("cluster", "Size")

  # Merge cluster means and sizes
  cluster_means_and_sizes <- merge(cluster_means, cluster_sizes_df, by = "cluster")

  # Store the calculated means and sizes
  cluster_means_and_sizes_list[[paste("MeansAndSizes_k", k)]] <- cluster_means_and_sizes

  # Print the means and sizes for each cluster for the current k
  cat("\nCluster Means and Sizes for k =", k, ":\n")
  print(cluster_means_and_sizes)
}

## 
## Cluster Means and Sizes for k = 3 :
##   cluster AMERICAN CULTURAL_IDENTITY KEEPSpan DISTINCT LEARNENG Size
## 1      1 3.473757      3.183011 3.261050 2.406077 3.243094  724
## 2      2 3.737692      3.428846 3.981154 2.701154 3.997308 2600

```

```

## 3      3 2.657084      3.591718 3.617385 2.981520 3.995209 1461
##
## Cluster Means and Sizes for k = 4 :
##   cluster AMERICAN CULTURAL_IDENTITY KEEPSPLAN DISTINCT LEARNENG Size
## 1      1 3.233256      3.411085 3.635104 2.711316 2.766744  433
## 2      2 3.880851      3.052482 3.327660 1.797163 3.988652  705
## 3      3 3.523810      3.158730 2.677249 2.997354 3.947090  378
## 4      4 3.256959      3.561946 3.996635 2.920465 4.000000 3269
##
## Cluster Means and Sizes for k = 5 :
##   cluster AMERICAN CULTURAL_IDENTITY KEEPSPLAN DISTINCT LEARNENG Size
## 1      1 3.149091      3.643636 4.000000 2.767273 2.752727  275
## 2      2 3.481132      2.858491 3.207547 2.339623 2.764151  106
## 3      3 3.301435      3.575279 3.955343 2.999681 3.971930 3135
## 4      4 3.344371      3.217715 3.973510 1.847682 4.000000  604
## 5      5 3.774436      3.022556 2.642105 2.393985 4.000000  665

```

Cluster Labeling and Profiles

```

# Step 1: Extract Cluster Assignments for k=4
cluster_assignments_k4 <- gmm_results[["GMM_k 4"]]\$classification

# Step 2: Add Cluster Assignments to the Dataset
# Assuming your dataset is named lns_subset
lns_subset\$cluster_assignment_k4 <- cluster_assignments_k4

# Verification Step 1: Confirm the column has been added
if("cluster_assignment_k4" %in% colnames(lns_subset)) {
  cat("Verification Passed: 'cluster_assignment_k4' column added.\n")
} else {
  cat("Verification Failed: Column not found.\n")
}

## Verification Passed: 'cluster_assignment_k4' column added.

# Verification Step 2: Check Cluster Sizes
cluster_sizes <- table(lns_subset\$cluster_assignment_k4)
cat("Cluster Sizes for k=4:\n")

## Cluster Sizes for k=4:
print(cluster_sizes)

##
##      1    2    3    4
##  433  705  378 3269

```

K-fold validation

```

# Generate folds ensuring indices are appropriate for the dataset
set.seed(2500)
folds <- createFolds(1:nrow(clustering_data_matrix), k = 5, list = TRUE)

# Function to perform GMM clustering and calculate silhouette score on test data
gmm_silhouette_cv <- function(train_indices, test_indices, data) {

```

```

train_data <- data[train_indices, ]
test_data <- data[test_indices, ]

# Fit GMM on the training data with k=4
gmm_fit <- Mclust(train_data, G = 4)

# Obtain cluster predictions for the test data
test_clusters <- predict(gmm_fit, newdata = test_data)$classification

# Calculate the silhouette score for the test data
silhouette_avg <- mean(silhouette(test_clusters, dist(test_data))[, "sil_width"])

return(silhouette_avg)
}

# Apply the cross-validation function across all folds
cv_silhouette_scores <- sapply(folds, function(indices) {
  gmm_silhouette_cv(indices, indices, clustering_data_matrix)
})

# Calculate mean and standard deviation of the silhouette scores across all folds
cv_mean_silhouette <- mean(cv_silhouette_scores)
cv_sd_silhouette <- sd(cv_silhouette_scores)

cat("Cross-Validation Mean Silhouette Score:", cv_mean_silhouette, "\n")

## Cross-Validation Mean Silhouette Score: 0.3566289
cat("Cross-Validation Silhouette Score SD:", cv_sd_silhouette, "\n")

## Cross-Validation Silhouette Score SD: 0.008071387

```

Bootstrapping (new)

Visualizations

Bootstrap

Parallel coordinates

```

# Extract the cluster means for k=4 using the correct key from the list
cluster_means_k4 <- cluster_means_and_sizes_list[["MeansAndSizes_k 4"]]

# Check the structure of cluster_means_k4 to ensure it's as expected
str(cluster_means_k4)

## 'data.frame':   4 obs. of  7 variables:
## $ cluster      : num  1 2 3 4
## $ AMERICAN    : num  3.23 3.88 3.52 3.26
## $ CULTURAL_IDENTITY: num  3.41 3.05 3.16 3.56
## $ KEEPSPLAN   : num  3.64 3.33 2.68 4
## $ DISTINCT    : num  2.71 1.8 3 2.92
## $ LEARNENG    : num  2.77 3.99 3.95 4
## $ Size         : int  433 705 378 3269

```

```

# If cluster_means_k4 is not NULL, print the first few rows to verify the content
if (!is.null(cluster_means_k4)) {
  print(head(cluster_means_k4))
} else {
  stop("Cluster means for k=4 not found in the list. Please check the list structure.")
}

##   cluster AMERICAN CULTURAL_IDENTITY KEEPSPLAN DISTINCT LEARNENG Size
## 1      1 3.233256          3.411085 3.635104 2.711316 2.766744 433
## 2      2 3.880851          3.052482 3.327660 1.797163 3.988652 705
## 3      3 3.523810          3.158730 2.677249 2.997354 3.947090 378
## 4      4 3.256959          3.561946 3.996635 2.920465 4.000000 3269

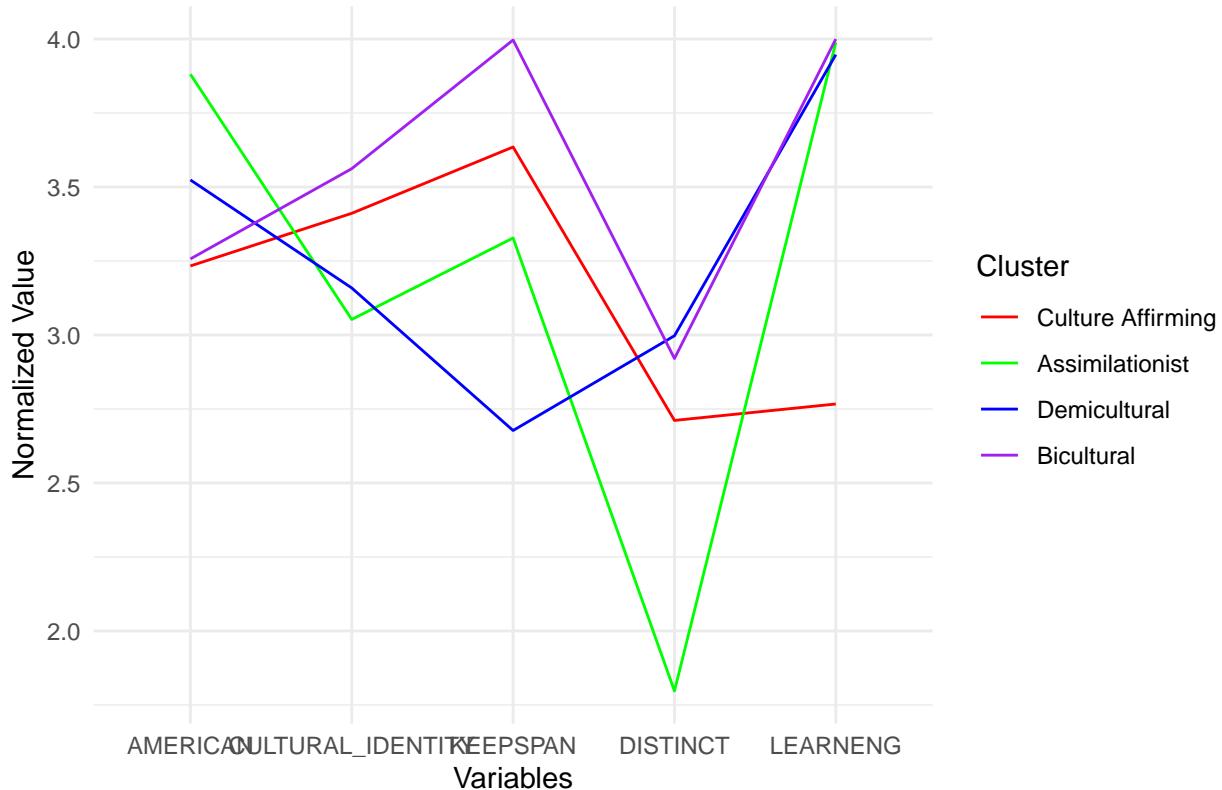
# Ensure the cluster column is correctly set as a factor
cluster_means_k4$cluster <- factor(cluster_means_k4$cluster, levels = c(1, 2, 3, 4),
                                     labels = c("Culture Affirming", "Assimilationist", "Demicultural", "Bicultural"))

# Create Parallel Coordinates Plot without the 'order' parameter
p <- ggparcoord(data = cluster_means_k4,
                  columns = 2:6, # Indices of the numeric columns to include in the plot
                  groupColumn = "cluster", # Name of the cluster column
                  scale = "globalminmax") +
  scale_color_manual(values = c("red", "green", "blue", "purple"),
                     labels = c("Culture Affirming", "Assimilationist", "Demicultural", "Bicultural"),
                     name = "Cluster") +
  labs(title = "Parallel Coordinates Plot of Cluster Means for k=4",
       x = "Variables",
       y = "Normalized Value") +
  theme_minimal() +
  theme(legend.position = "right")

# Display plot
print(p)

```

Parallel Coordinates Plot of Cluster Means for k=4



Radar Chart

```
# Creating a data frame for the radar chart with minimum and maximum values
radar_data <- data.frame(
  cluster = factor(c("Culture Affirming", "Assimilationist", "Demicultural", "Bicultural")),
  AMERICAN = c(3.23, 3.88, 3.52, 3.26),
  CULTURAL_IDENTITY = c(3.41, 3.05, 3.16, 3.56),
  KEEPSPLAN = c(3.64, 3.33, 2.68, 4),
  DISTINCT = c(2.71, 1.8, 3, 2.92),
  LEARNENG = c(2.77, 3.99, 3.95, 4)
)

# Define the maximum and minimum values for each variable
max_values <- apply(radar_data[, -1], 2, max)
min_values <- apply(radar_data[, -1], 2, min)

# Normalize the radar_data to the scale of the plot
radar_data_normalized <- as.data.frame(mapply(function(x, min, max) (x - min) / (max - min),
                                               radar_data[, -1], min_values, max_values))

# Creating the radar chart data
radar_chart_data <- rbind(
  min = rep(0, times=ncol(radar_data_normalized)),
  max = rep(1, times=ncol(radar_data_normalized)),
  radar_data_normalized
)
```

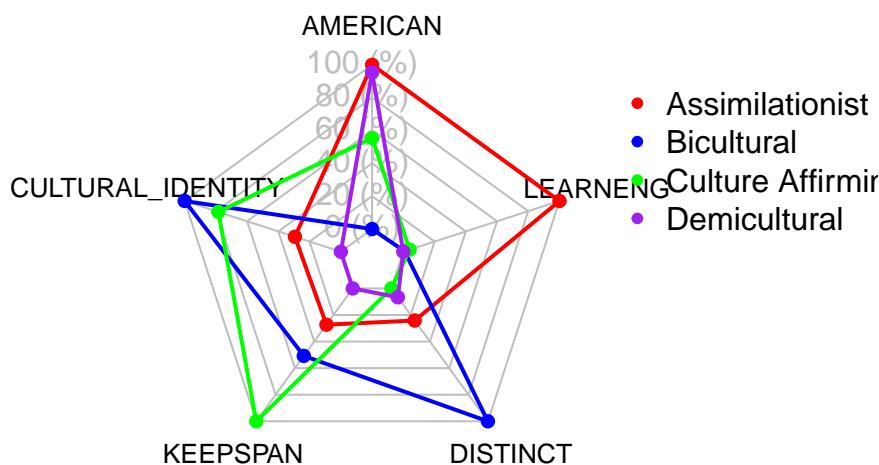
```

# Adding cluster labels as row names
rownames(radar_chart_data)[3:6] <- as.character(radar_data$cluster)

# Plot the radar chart
radarchart(radar_chart_data, axistype = 1,
           seg = 5,
           pcol = c("red", "blue", "green", "purple"),
           plty = 1,
           plwd = 2,
           cglcol="grey", cglty=1, axislabcol="grey",
           vlcex=0.8)

# Add a legend outside the plot
legend(x=1.2, y=1, legend = levels(radar_data$cluster), bty = "n", pch=20, col=c("red", "blue", "green"))

```



Save the clusters in an RDA file

```

# Ensure the cluster_assignment_k4 variable is a factor
lns_subset$cluster_assignment_k4 <- factor(lns_subset$cluster_assignment_k4)

# Define the labels for each cluster
cluster_labels <- c('Culture Affirming', 'Assimilationist', 'Demicultural', 'Bicultural')

# Assign names to the levels of the clusters
levels(lns_subset$cluster_assignment_k4) <- cluster_labels

# Rename cluster_assignment_k4 to 'clusters'
names(lns_subset)[names(lns_subset) == 'cluster_assignment_k4'] <- 'clusters'

# Verify the renaming and releveling
str(lns_subset$clusters)

## Factor w/ 4 levels "Culture Affirming",...: 4 4 4 1 4 4 4 4 4 3 ...
# Check the size of each cluster to ensure it aligns with the given sizes
print(table(lns_subset$clusters))

##
## Culture Affirming    Assimilationist      Demicultural        Bicultural

```

```
##          433          705          378         3269
# Save the dataset with the new cluster labels
save(lns_subset, file = "~/Desktop/Dissertation Analysis/Dissertation Analysis/Data/Processed/lns_clust
```

End of Script